

Practical Numerical Optimization with Scipy, Estimagic and JAXopt

Scipy Conference 2022

Janos Gabler & Tim Mensinger

University of Bonn

About Us



- Website: janosg.com
- GitHub: [janosg](https://github.com/janosg)
- Original author of estimagic
- Submitted PhD thesis, looking for interesting jobs soon



- Website: tmensinger.com
- GitHub: [timmens](https://github.com/timmens)
- estimagic core contributor
- PhD student in Econ, University of Bonn

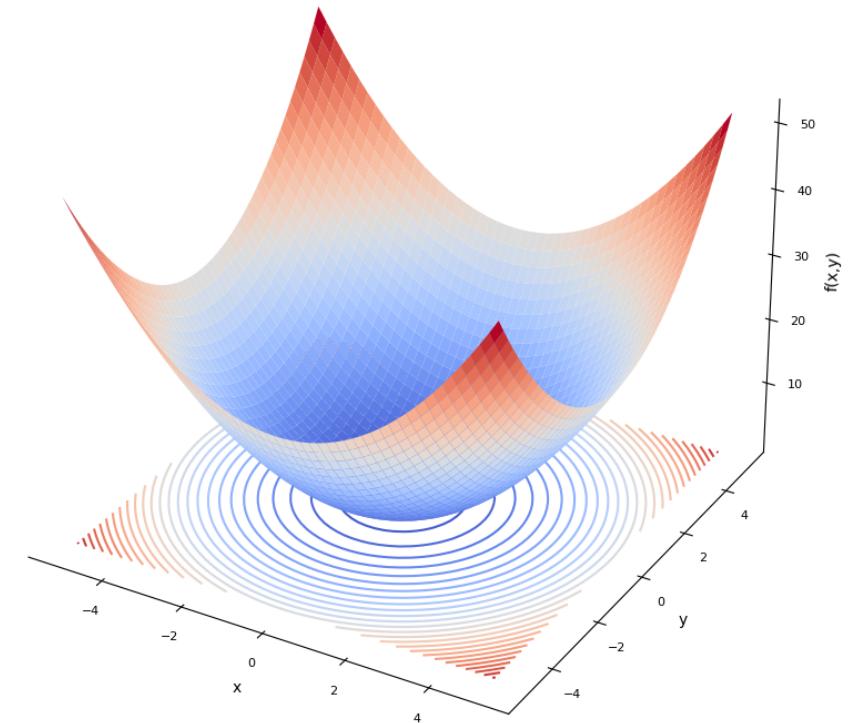
Index

1. What is numerical optimization
2. Introduction to `scipy.optimize`
3. Introduction to `estimagic`
4. Choosing algorithms
5. Advanced `estimagic`
6. Jax and Jaxopt

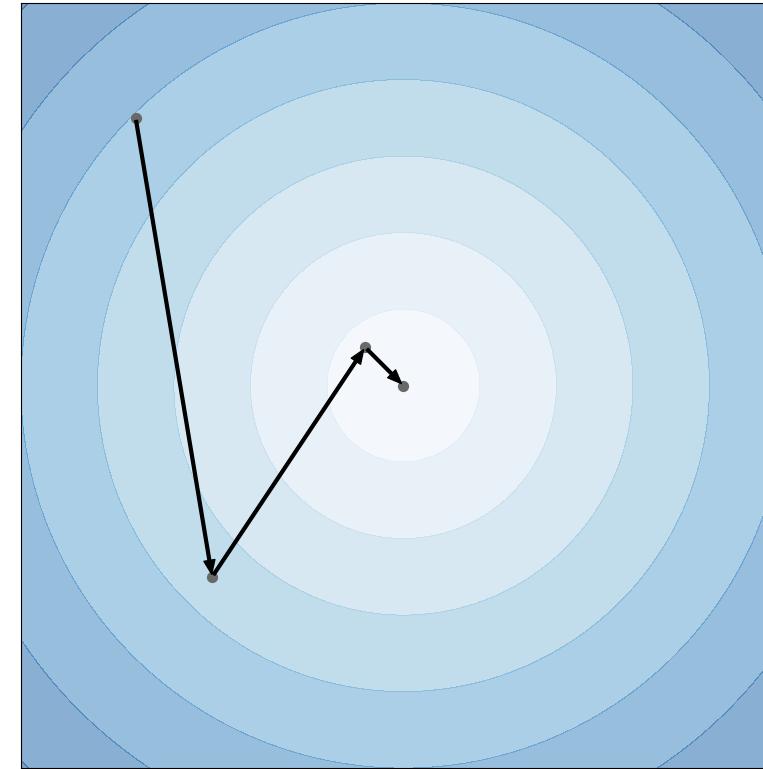
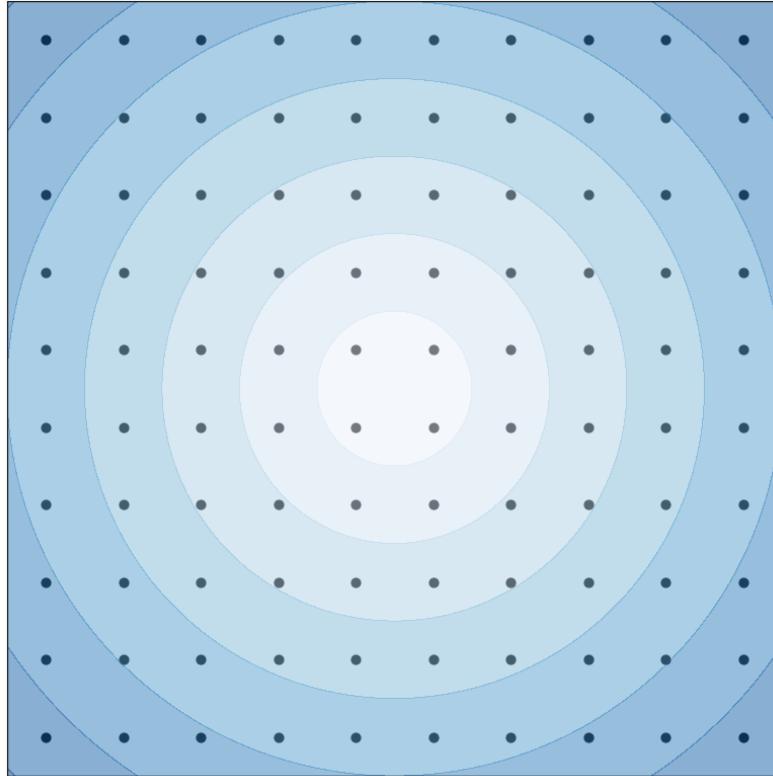
What is numerical optimization

Example problem

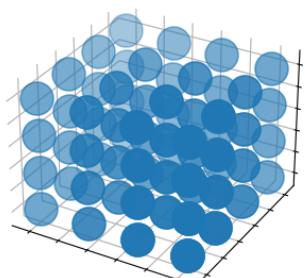
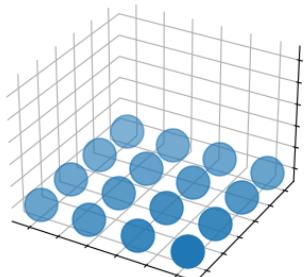
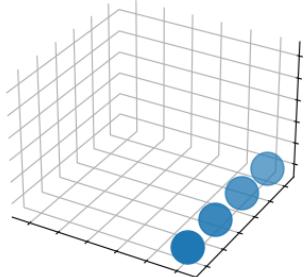
- Parameters x_1, x_2
- Criterion $f(x_1, x_2) = x_1^2 + x_2^2$
- Want: $x_1^*, x_2^* = \operatorname{argmin} f(x_1, x_2)$
- Possible extensions:
 - Constraints
 - Bounds
- Optimum at $(0, 0)$ with function value 0



Brute force vs. smarter algorithm



Complexity of brute force



Number of Dimensions	Runtime (1 ms per evaluation, 100 points per dimension)
1	100 ms
2	10 s
3	16 min
4	27 hours
5	16 weeks
6	30 years

In this talk

- Nonlinear optimization with
- Continuous parameters
- (Non)linear constraints
- Global optimization
- Diagnostics and strategies for difficult problems

Not covered

- Linear programming
- Mixed integer programming
- Stochastic gradient descent
- Fitting neural networks

Introduction to `scipy.optimize`

Solve example problem with `scipy.optimize`

```
>>> import numpy as np
>>> from scipy.optimize import minimize

>>> def sphere(x):
...     return np.sum(x ** 2)

>>> x0 = np.ones(2)
>>> res = minimize(f, x0)
>>> res.fun
0.0
>>> res.x
array([0.0, 0.0])
```

Features of `scipy.optimize`

- `minimize` as unified interface to 14 local optimizers
 - some support bounds
 - some support constraints
- Parameters are 1d arrays
- Maximization is done by minimizing $-f(x)$
- Different interfaces for:
 - global optimization
 - nonlinear least-squares

Practice Session 1: First optimization with `scipy.optimize` (15 min)

Shortcomings of `scipy.optimize`

- Very few algorithms
- No parallelization
- Maximization via sign flipping
- No diagnostics tools
- No feedback before termination or in case of crash
- No built-in multistart, benchmarking, scaling, or logging
- Parameters are 1d numpy arrays

Examples from real projects I

```
def parse_parameters(x):
    """Parse the parameter vector into quantities we need."""
    num_types = int(len(x[54:]) / 6) + 1
    params = {
        'delta': x[0:1],
        'level': x[1:2],
        'coeffs_common': x[2:4],
        'coeffs_a': x[4:19],
        'coeffs_b': x[19:34],
        'coeffs_edu': x[34:41],
        'coeffs_home': x[41:44],
        'type_shares': x[44:44 + (num_types - 1) * 2],
        'type_shifts': x[44 + (num_types - 1) * 2:]
    }
    return params
```

Examples from real projects II

```
>>> scipy.optimize.minimize(func, x0)
-----
LinAlgError                                     Traceback (most recent call last)
<ipython-input-17-7459e5b4d8d4> in <module>
----> 1 scipy.optimize.minimize(func, x0)

 95
 96 def _raise_linalgerror_singular(err, flag):
---> 97     raise LinAlgError("Singular matrix")
 98

LinAlgError: Singular matrix
```

- After 5 hours and with no additional information

Introduction to estimagic

What is estimagic?

- Library for difficult numerical optimization
- Tools for nonlinear estimation and inference
- Wraps many other optimizer libraries:
 - Scipy, Nlopt, TAO, Pygmo, ...
- Harmonized interface
- A lot of additional functionality

You can use it like scipy

```
>>> import estimagic as em

>>> def sphere(x):
...     return np.sum(x ** 2)

>>> res = em.minimize(
...     criterion=sphere,
...     params=np.arange(5),
...     algorithm="scipy_lbfgsb",
... )

>>> res.params
array([ 0., -0., -0., -0., -0.])
```

- But there is no default algorithm

Params can be (almost) anything

```
>>> def dict_sphere(x):
...     out = (
...         x["a"] ** 2 + x["b"] ** 2 + (x["c"] ** 2).sum()
...     )
...     return out

>>> res = em.minimize(
...     criterion=dict_sphere,
...     params={"a": 0, "b": 1, "c": pd.Series([2, 3, 4])},
...     algorithm="scipy_neldermead",
... )

>>> res.params
{'a': 0.,
 'b': 0.,
 'c': 0    0.
      1    0.
      2    0.
      dtype: float64}
```

- numpy arrays
- pd.Series, pd.DataFrame
- scalars
- (Nested) lists, dicts and tuples thereof
- Special case: DataFrame with columns "value", "lower_bound" and "upper_bound"

Informative OptimizeResult

```
>>> res
Minimize with 5 free parameters terminated successfully after 805 criterion evaluations and 507 iterations.
```

The value of criterion improved from 30.0 to 1.6760003634613059e-16.

The `scipy_neldermead` algorithm reported: Optimization terminated successfully.

Independent of the convergence criteria used by `scipy_neldermead`, the strength of convergence can be assessed by the following criteria:

	one_step	five_steps
relative_criterion_change	1.968e-15***	2.746e-15***
relative_params_change	9.834e-08*	1.525e-07*
absolute_criterion_change	1.968e-16***	2.746e-16***
absolute_params_change	9.834e-09**	1.525e-08*

(***: change <= 1e-10, **: change <= 1e-8, *: change <= 1e-5. Change refers to a change between accepted steps. The first column only considers the last step. The second column considers the last five steps.)

Access OptimizeResult's attributes

```
>>> res.criterion  
.0  
>>> res.n_criterion_evaluations  
805  
>>> res.success  
True  
>>> res.message  
'Optimization terminated successfully.'  
>>> res.history.keys()  
dict_keys(['params', 'criterion', 'runtime'])
```

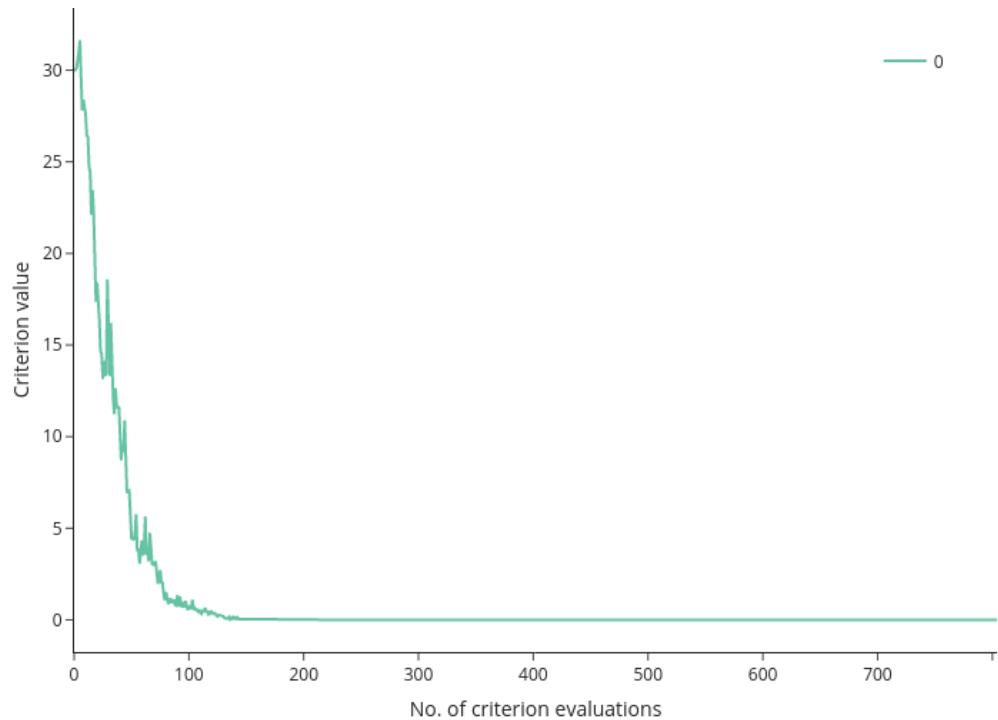
Logging and Dashboard

```
>>> res = em.minimize(  
...     criterion=sphere,  
...     params=np.arange(5),  
...     algorithm="scipy_lbfgsb",  
...     logging="my_log.db",  
...     log_options={  
...         "if_database_exists": "replace"  
...     },  
... )  
  
>>> from estimagic import OptimizeLogReader  
  
>>> reader = OptimizeLogReader("my_log.db")  
>>> reader.read_history().keys()  
dict_keys(['params', 'criterion', 'runtime'])  
  
>>> reader.read_iteration(1)["params"]  
array([0., 0.817, 1.635, 2.452, 3.27])
```

- Persistent log of parameters and criterion values
- Thread safe sqlite database
- No data loss, even if computer is unplugged
- Can be read while optimization is running
- Provides data for dashboard

Criterion plot

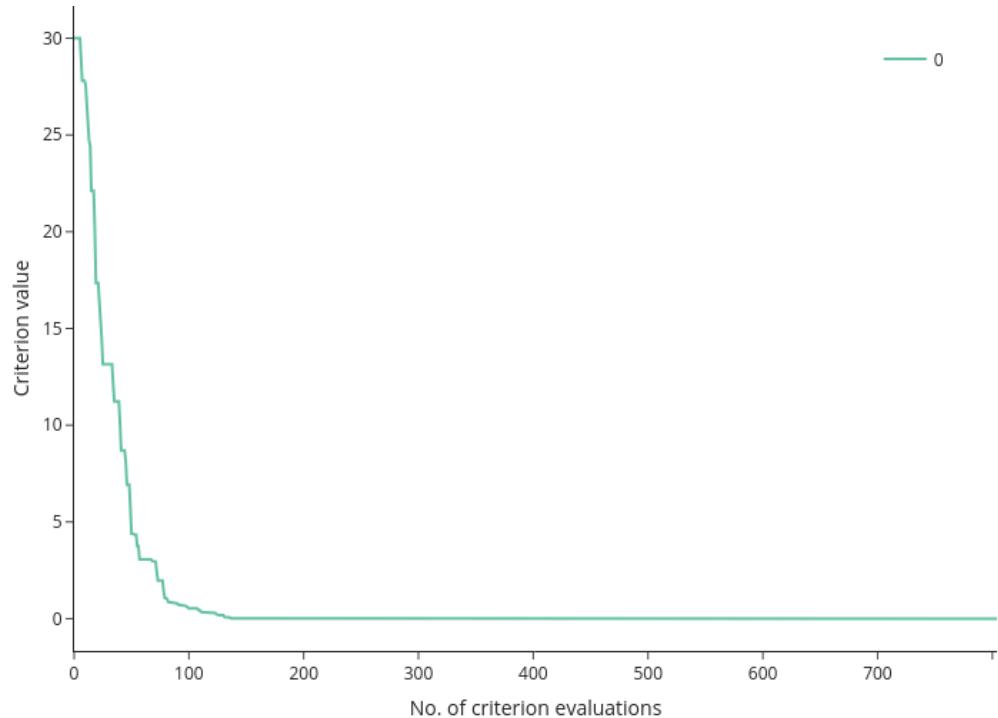
```
em.criterion_plot(res)
```



- First argument can be:
 - OptimizeResult
 - path to log file
 - list or dict thereof
- Dictionary keys are used for legend

Criterion plot (2)

```
em.criterion_plot(res, monotone=True)
```

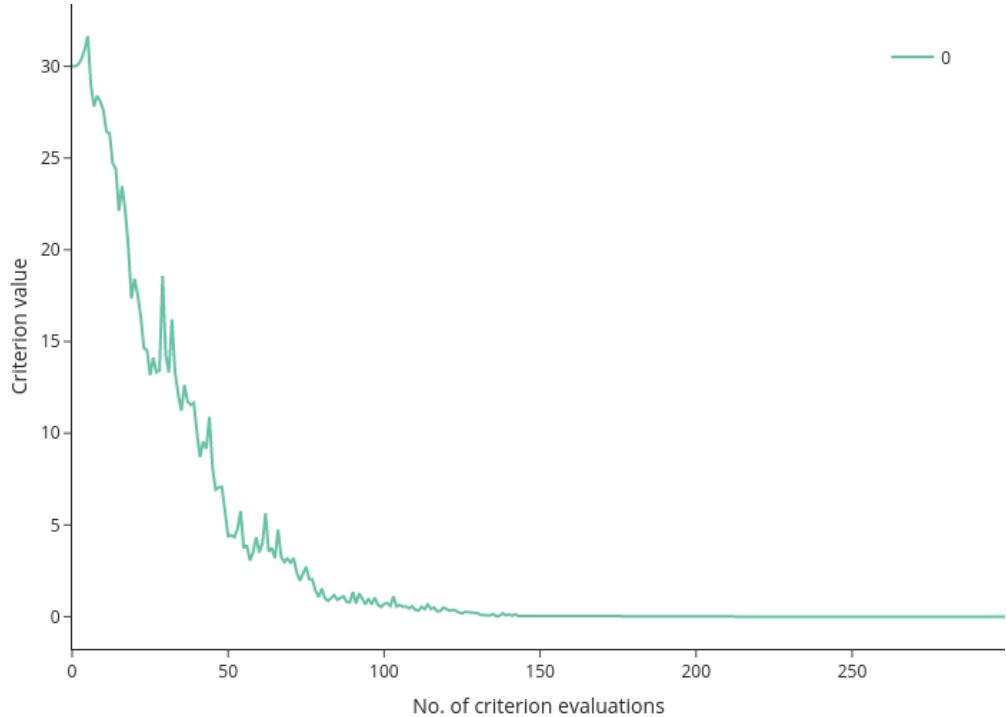


- `monotone=True` shows the current best value
- useful if there are extreme values in history

Criterion plot (3)

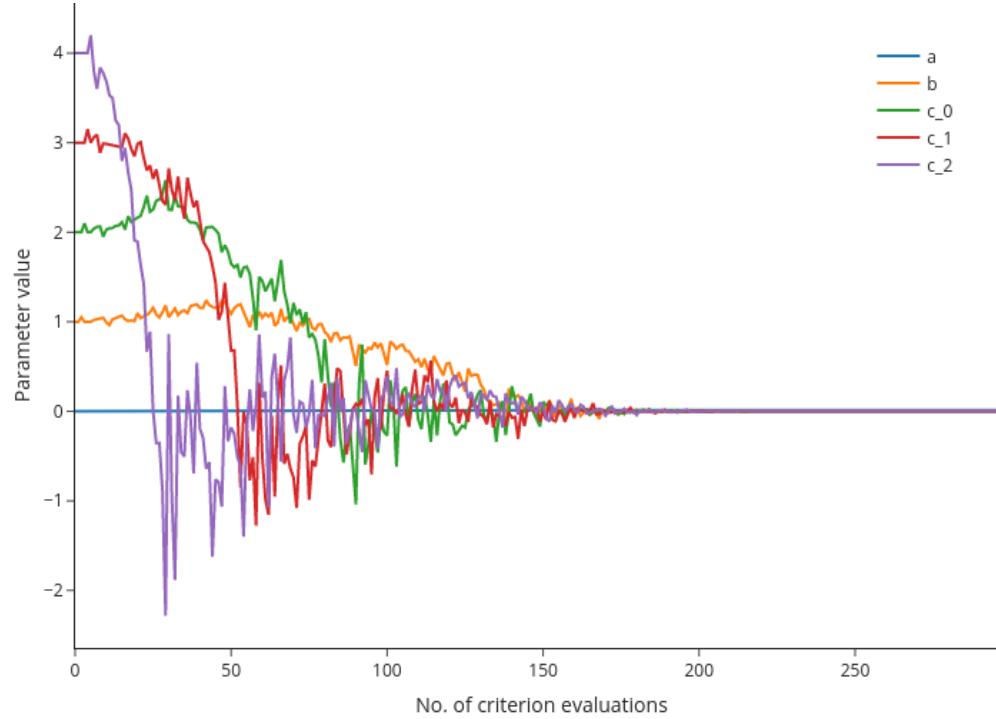
```
em.criterion_plot(res, max_evaluations=300)
```

- `max_evaluations` sets upper limit of x-axis



Params plot

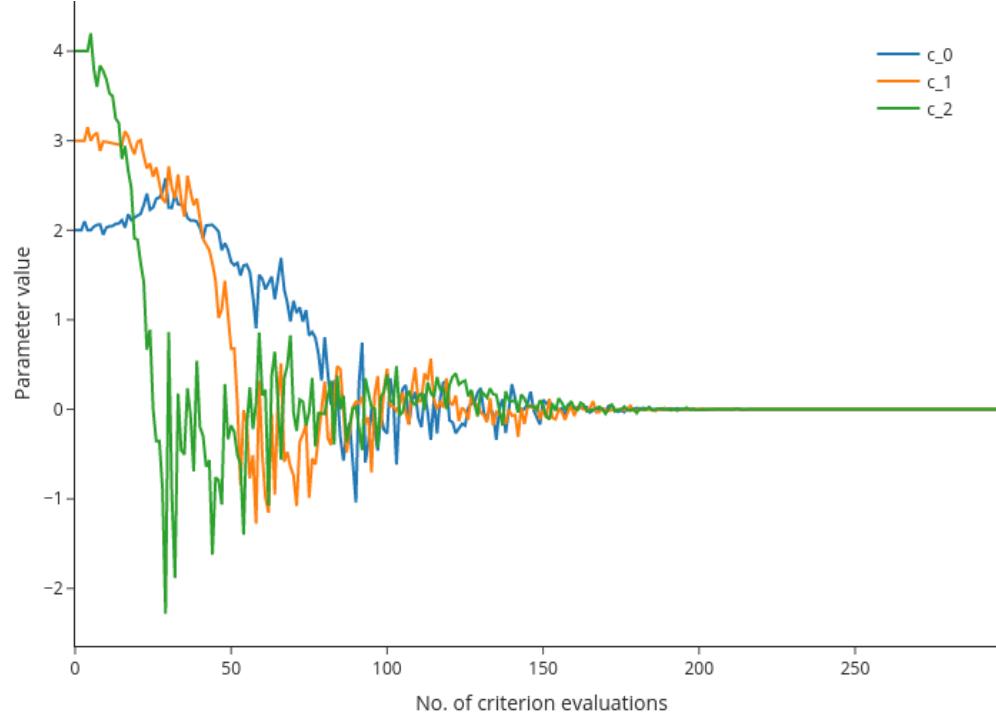
```
params = {  
    "a": 0,  
    "b": 1,  
    "c": pd.Series([2, 3, 4])  
}  
em.params_plot(  
    res,  
    max_evaluations=300,  
)
```



- Similar options as `criterion_plot`

Params plot (2)

```
params = {
    "a": 0,
    "b": 1,
    "c": pd.Series([2, 3, 4])
}
em.params_plot(
    res,
    max_evaluations=300,
    selector=lambda x: x["c"],
)
```



- selector is a function returning a subset of params

Unified interface to (collections of) algorithms

- **Nlopt**: Local and global algorithms by Stephen Johnson (MIT)
- **Pygmo**: Global algorithms by Francesco Biscani and Dario Izzo (ESA)
- **fides**: Pure Python algorithm by Fabian Fröhlich (Harvard)
- **TAO**: Toolkit for advanced optimization (Argonne national lab)
- **Cyipopt**: Python bindings to **IPOPT** by Andreas Wächter (Northwestern)
- **NAG**: Numerical algorithms group (Oxford)
- **estimagic**: A few algorithms we could not find elsewhere
- **List not exhaustive and will add more soon (it is quite easy)**

Use constraints with any optimizer

```
>>> res = em.minimize(  
...     criterion=sphere,  
...     params=np.array([0.1, 0.5, 0.4, 4, 5]),  
...     algorithm="scipy_lbfgsb",  
...     constraints=[{  
...         "loc": [0, 1, 2],  
...         "type": "probability"  
...     }],  
... )  
  
>>> res.params  
array([0.33334, 0.33333, 0.33333, -0., 0.])
```

- lbfgsb is unconstrained
- estimagic transforms constrained problems into unconstrained ones
- Supported constraints:
 - linear
 - probability
 - covariance
 - ...

Closed-form or parallel numerical derivatives

```
>>> def sphere_gradient(params):
...     return 2 * params

>>> em.minimize(
...     criterion=sphere,
...     params=np.arange(5),
...     algorithm="scipy_lbfgsb",
...     derivative=sphere_gradient,
... )

>>> em.minimize(
...     criterion=sphere,
...     params=np.arange(5),
...     algorithm="scipy_lbfgsb",
...     numdiff_options={"n_cores": 6},
... )
```

- You can provide derivatives
- Otherwise, estimagic calculates them numerically
- Parallelization on (up to) as many cores as parameters

There is maximize

```
>>> def upside_down_sphere(params):
...     return -params @ params

>>> res = em.maximize(
...     criterion=upside_down_sphere,
...     params=np.arange(5),
...     algorithm="scipy_lbfgs",
... )
>>> res.params
array([ 0.,  0.,  0.,  0.,  0.])
```

Multistart framework

```
>>> res = em.minimize(  
...     criterion=sphere,  
...     params=np.arange(5),  
...     algorithm="scipy_neldermead",  
...     soft_lower_bounds=np.full(5, -5),  
...     soft_upper_bounds=np.full(5, 15),  
...     multistart=True,  
...     multistart_options={  
...         "convergence.max_discoveries": 5,  
...         "n_samples": 1000  
...     },  
... )  
>>> res.params  
array([0., 0., 0., 0., 0.])
```

- Turn local optimizers global
- Inspired by [tiktak algorithm](#)
- **Exploration** on random sample
- Local optimizations from best points
- Use any optimizer

Exploit structure of F

```
>>> def general_sphere(params):
...     contribs = params ** 2
...     out = {
...         "root_contributions": params,
...         "contributions": contribs,
...         "value": contribs.sum(),
...     }
...     return out

>>> res = em.minimize(
...     criterion=general_sphere,
...     params=np.arange(5),
...     algorithm="pounders",
... )
>>> res.params
array([0., 0., 0., 0., 0.])
```

- Common structures
 - $F(x) = \sum_i f_i(x)^2$ (least squares)
 - $F(x) = \sum_i f_i(x)$ (e.g. log-likelihood)
- Huge speed-ups
- Increased robustness

Harmonized as much as possible but not more

```
>>> algo_options = {  
...     "convergence.relative_criterion_tolerance": 1e-9,  
...     "stopping.max_iterations": 100_000,  
...     "trustregion.initial_radius": 10.0,  
... }  
  
>>> res = em.minimize(  
...     criterion=sphere,  
...     params=np.arange(5),  
...     algorithm="nag_pybobyqa",  
...     algo_options=algo_options,  
... )  
>>> res.params  
array([0., 0., 0., 0., 0.])
```

- Same options have same name
- Different options have different names
 - e.g. not one tol
- Ignore irrelevant options

Well documented

estimagic

Getting Started How-to Guides Explanations API Development Optimizers

Previous topic
[estimagic](#)

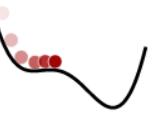
Next topic
[Installation](#)

Quick search

[Go](#)

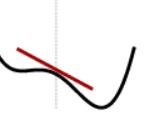
Getting Started

This section contains quickstart guides for new estimagic users. It can also serve as a reference for more experienced users.



Optimization

Learn numerical optimization with estimagic



Differentiation

Learn numerical differentiation with estimagic



Estimation

Learn maximum likelihood and methods of simulated moments estimation with estimagic

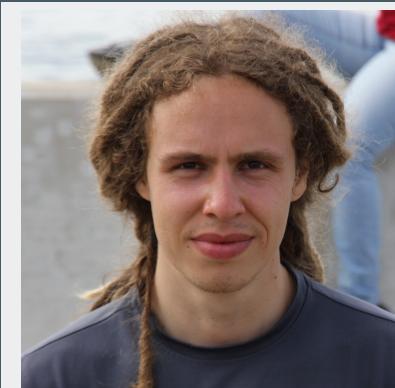


Installation

Installation instructions for estimagic and optional dependencies

- [estimagic.readthedocs.io](#)
- Quickstart
- Detailed How-To guides
- Explanations
- API

The estimagic Team



Janos



Tim



Klara



Sebastian



Tobias



Hans-Martin

Thanks to

- All [contributors](#) to estimagic
- [Kenneth Judd](#) for feedback and funding of a research visit
- [Gregor Reich](#) for feedback
- All authors of the amazing algorithms we are wrapping
- The University of Bonn and [TRA-1 Modelling](#) for funding
- [Collaborative Research Center Transregio 224](#) for funding

Break (5 min)

Practice Session 2: Convert previous example to estimagic (15 min)

Choosing algorithms

Relevant problem properties

- **Smoothness:** Differentiable? Kinks? Discontinuities? Stochastic?
- **Convexity:** Are there local optima?
- **Size:** 2 parameters? 10? 100? 1000? More?
- **Constraints:** Bounds? Linear constraints? Nonlinear constraints?
- **Structure:** Nonlinear least-squares, Log-likelihood function
- **Goal:** Do you need a global solution? How precise?
- Properties guide selection but experimentation is important

`scipy_lbfgsb`

- Limited memory BFGS
- BFGS: Approximate hessians from multiple gradients
- Criterion must be differentiable
- Scales to a few thousand parameters
- Beats other BFGS implementations in many benchmarks
- Low overhead

fides

- Derivative based trust-region algorithm
- Developed by Fabian Fröhlich as a Python package
- Many advanced options to customize the optimization!
- Criterion must be differentiable
- Good solution if `scipy_lbfgsb` is too aggressive

nlopt_bobyqa , nag_pybobyqa

- **Bound Optimization by Quadratic Approximation**
- Derivative free trust region algorithm
- nlopt has less overhead
- nag has options to deal with noise
- Good for non-differentiable not too noisy functions
- Slower than derivative based methods but faster than neldermead

scipy_neldermead , nlopt_neldermead

- Popular direct search method
- scipy does not support bounds
- nlopt requires fewer criterion evaluations in most benchmarks
- Almost never the best choice but sometimes not the worst

`scipy_ls_lm`, `scipy_ls_trf`

- Derivative based optimizers for nonlinear least squares
- Criterion needs structure: $F(x) = \sum_i f_i(x)^2$
- In estimagic, criterion returns a dict:

```
def sphere_ls(x):
    # x are the least squares residuals in the sphere function
    return {"root_contributions": x, "value": x @ x}
```

- `scipy_ls_lm` is better for small problems without bounds
- `scipy_ls_trf` is better for problems with many parameters

nag_dfols , pounders

- Derivative free trust region method for nonlinear least-squares
- Both beat bobyqa for least-squares problems!
- nag_dfols is usually fastest
- nag_dfols has advanced options to deal with noise
- pounders can do criterion evaluations in parallel

ipopt

- Interior point optimizer for problems with nonlinear constraints
- Probably best open source optimizer for nonlinear constraints
- We wrap it via `cyipopt`

Practice Session 3: Play with algorithm and algo_options (20 min)

What is benchmarking

- Compare algorithms on functions with known optimum
- Mirror problems you actually want to solve
 - similar number of parameters
 - similar w.r.t. differentiability or noise
- Benchmark functions should be fast!

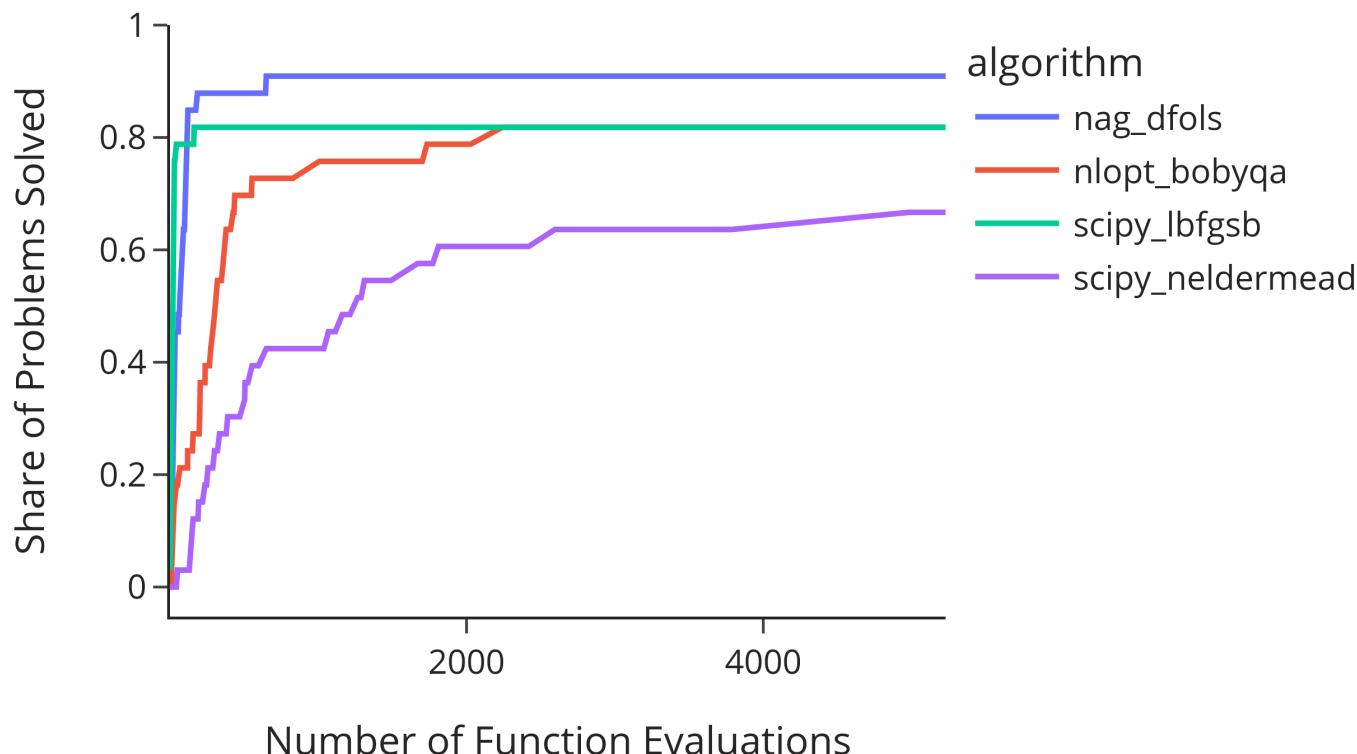
Running benchmarks in estimagic

```
problems = em.get_benchmark_problems(  
    "estimagic",  
)  
optimizers = [  
    "scipy_lbfgsb",  
    "nag_dfols",  
    "nlopt_bobyqa",  
    "scipy_neldermead",  
]  
results = em.run_benchmark(  
    problems=problems,  
    optimize_options=optimizers,  
    n_cores=4,  
    max_criterion_evaluations=1000,  
)
```

1. Create a set of test problems:
 - dict with functions, start values and correct solutions
2. Define optimize_options
 - Optimizers to try
 - Optionally: keyword arguments for minimize
3. Get benchmark results
 - Dict with histories and solutions

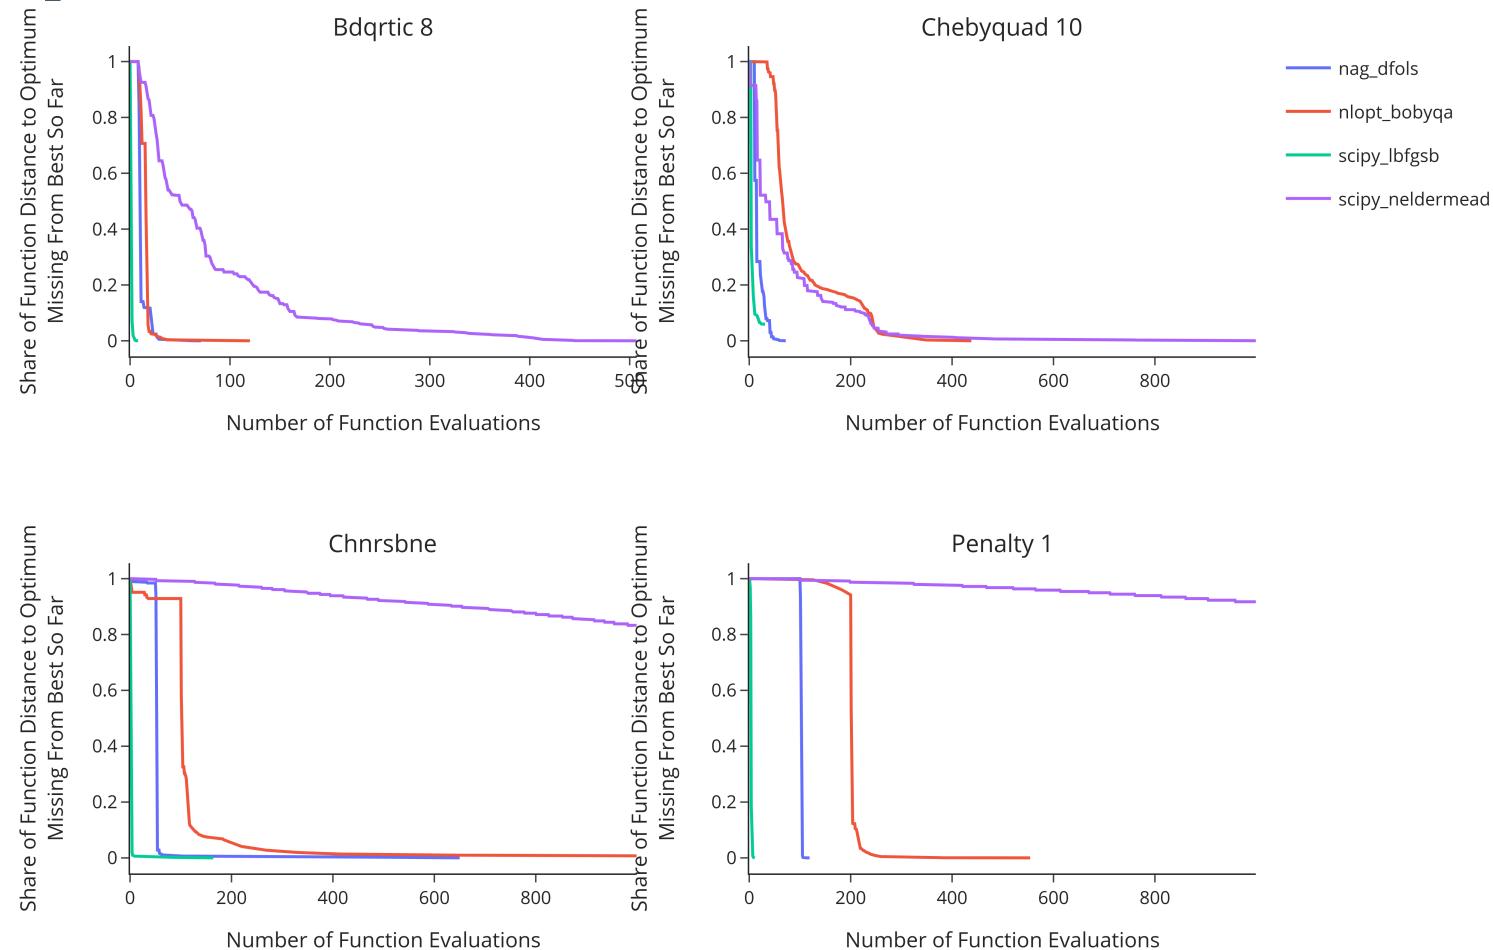
Profile plots

```
em.profile_plot(problems, results)
```



Convergence plots

```
subset = [
    "chebyquad_10",
    "chnrsbne",
    "penalty_1",
    "bdqrtic_8",
]
em.convergence_plot(
    problems,
    results,
    problem_subset=subset,
)
```



Why not look at runtime

- Benchmark functions are very fast (microseconds)
-> Runtime dominated by algorithm overhead
- Most real functions are slow (milliseconds, seconds, minutes, ...)
-> Runtime dominated by number of evaluations
- Exceptions will be discussed later (jaxopt)

Advanced options

```
problems = em.get_benchmark_problems(  
    name="example",  
    additive_noise=True,  
    additive_noise_options={  
        "distribution": "normal",  
        "std": 0.2,  
    },  
    scaling=True,  
    scaling_options={  
        "min_scale": 0.1,  
        "max_scale": 1000,  
    }  
)
```

- Add noise and scaling issues
- Choices:
 - estimagic: small and large problems
 - more_wild: small least squares problems
 - example: subset of more_wild

Practice Session 4: Benchmarking optimizers (10 min)

Break (10 min)

Advanced estimagic

Constraints

Terminology of constraints in estimagic

- Constraints are conditions on parameters
- Can make problems simpler or more difficult
- bounds: $\min_x f(x)$ s.t. $l \leq x \leq u$
 - handled by most algorithms
- estimagic constraints:
 - handled by estimagic via reparametrization
- nonlinear constraints: $\min_x f(x)$ s.t. $c_1(x) = 0, c_2(x) \geq 0$
 - handled by some algorithms
 - can be violated during optimization

How to specify bounds for array params

```
>>> def sphere(x):
...     return x @ x

>>> res = em.minimize(
...     criterion=sphere,
...     params=np.arange(3) + 1,
...     lower_bounds=np.ones(3),
...     algorithm="scipy_lbfgsb",
... )
>>> res.params
array([1., 1., 1.])
```

- Specify lower_bounds and upper_bounds
- Use np.inf or -np.inf to represent no bounds

How to specify bounds for DataFrame params

```
>>> params = pd.DataFrame({  
...     "value": [1, 2, 3],  
...     "lower_bound": [1, 1, 1],  
...     "upper_bound": [3, 3, 3],  
... },  
...     index=[ "a", "b", "c" ],  
... )  
  
>>> def criterion(p):  
...     return (p[ "value" ] ** 2).sum()  
  
>>> em.minimize(criterion, params, algorithm="scipy_lbfgsb")
```

How to specify bounds for pytree params

```
params = {"x": np.arange(3), "intercept": 3}

def criterion(p):
    return p["x"] @ p["x"] + p["intercept"]

res = em.minimize(
    criterion,
    params=params,
    algorithm="scipy_lbfgsb",
    lower_bounds={"intercept": -2},
)
```

- Enough to specify the subset of params that actually has bounds
- We try to match your bounds with params
- Raise `InvalidBoundsError` in case of ambiguity

Reparametrization example

- Example: $f(x_1, x_2) = \sqrt{x_2 - x_1} + x_2^2$
- Only defined for $x_1 \leq x_2$
- Solve: $\min f(x_1, x_2)$ s.t. $x_1 \leq x_2$
- Not a simple bound!
- Reparametrization approach:
 - Define $\tilde{x}_2 = x_2 - x_1$ and $\tilde{f}(x_1, \tilde{x}_2) = \sqrt{\tilde{x}_2} + (x_1 + \tilde{x}_2)^2$
 - Solve: $\min \tilde{f}(x_1, \tilde{x}_2)$ s.t. $\tilde{x}_2 \geq 0$
 - Translate solution back into x_1 and x_2

Which constraints can be handled via reparametrization?

- Fixing parameters (simple but useful)
- Find valid covariance and correlation matrices
- Find valid probabilities
- Linear constraints (as long as there are not too many)
 - $\min f(x) \text{ s.t. } A_1x = 0 \text{ and } A_2x \leq 0$
- **Guaranteed to be fulfilled during optimization**

Do not try at home

- Easy to make mistakes when implementing this
 - forget to transform start parameters
 - forget to translate back
 - forget to adjust derivative
 - confuse directions
 - use non-differentiable transformations
- **Estimagic does reparametrizations for you!**
- **Completely hides transformed x**

Fixing parameters

```
>>> def criterion(params):
...     offset = np.linspace(1, 0, len(params))
...     x = params - offset
...     return x @ x

unconstrained_optimum = [1, 0.8, 0.6, 0.4, 0.2, 0]

>>> res = em.minimize(
...     criterion=criterion,
...     params=np.array([2.5, 1, 1, 1, 1, -2.5]),
...     algorithm="scipy_lbfgsb",
...     constraints={"loc": [0, 5], "type": "fixed"},
... )
>>> res.params
array([ 2.5,  0.8,  0.6,  0.4,  0.2, -2.5])
```

- loc selects location 0 and 5 of the parameters
- type states that they are fixed
- Other selection methods will be explained later

Linear constraints

```
>>> res = em.minimize(  
...      criterion=criterion,  
...      params=np.ones(6),  
...      algorithm="scipy_lbfgsb",  
...      constraints=  
...          {"loc": [0, 1, 2, 3],  
...           "type": "linear",  
...           "lower_bound": 0.95,  
...           "weights": 0.25,  
...         },  
...       )  
>>> res.params  
array([ 1.25,  1.05,  0.85,  0.65,  0.2 , -0.])
```

- Imposes that average of first 4 parameters is larger than 0.95
- Weights can be scalars or same length as selected parameters
- Use "value" instead of "lower_bound" for linear equality constraint

Constraints have to hold

```
>>> em.minimize(  
...     criterion=sphere,  
...     params=np.array([1, 2, 3, 4, 5]),  
...     algorithm="scipy_lbfgsb",  
...     constraints={"loc": [0, 1, 2], "type": "probability"},  
... )
```

InvalidParamsError: A constraint of type 'probability' is not fulfilled in params, please make sure that it holds for the starting values. The problem arose because: Probabilities do not sum to 1. The names of the involved parameters are: ['0', '1', '2'] The relevant constraint is:
{'loc': [0, 1, 2], 'type': 'probability', 'index': array([0, 1, 2])}.

Nonlinear constraints

```
>>> res = em.minimize(  
...     criterion=criterion,  
...     params=np.ones(6),  
...     algorithm="scipy_slsqp",  
...     constraints={  
...         "type": "nonlinear",  
...         "loc": [0, 1, 2, 3, 4],  
...         "func": lambda x: np.prod(x),  
...         "value": 1.0,  
...     },  
... )  
>>> res.params  
array([1.31, 1.16, 1.01, 0.87, 0.75, -0.])
```

- Restrict the product of first five parameters to 1
- Only works with some optimizers
- func can be an arbitrary python function of params that returns a number, array or pytree
- Use "lower_bound" and "upper_bound" instead of "value" for inequality constraints

Parameter selection methods

- "loc" can be replaced other things
- If params is a DataFrame with "value" column
 - "query" : An arbitrary query string that selects the relevant rows
 - "loc" : Will be passed to `DataFrame.loc`
- Always
 - "selector" : A function that takes params as argument and returns a subset of params
- More in the [documentation](#)

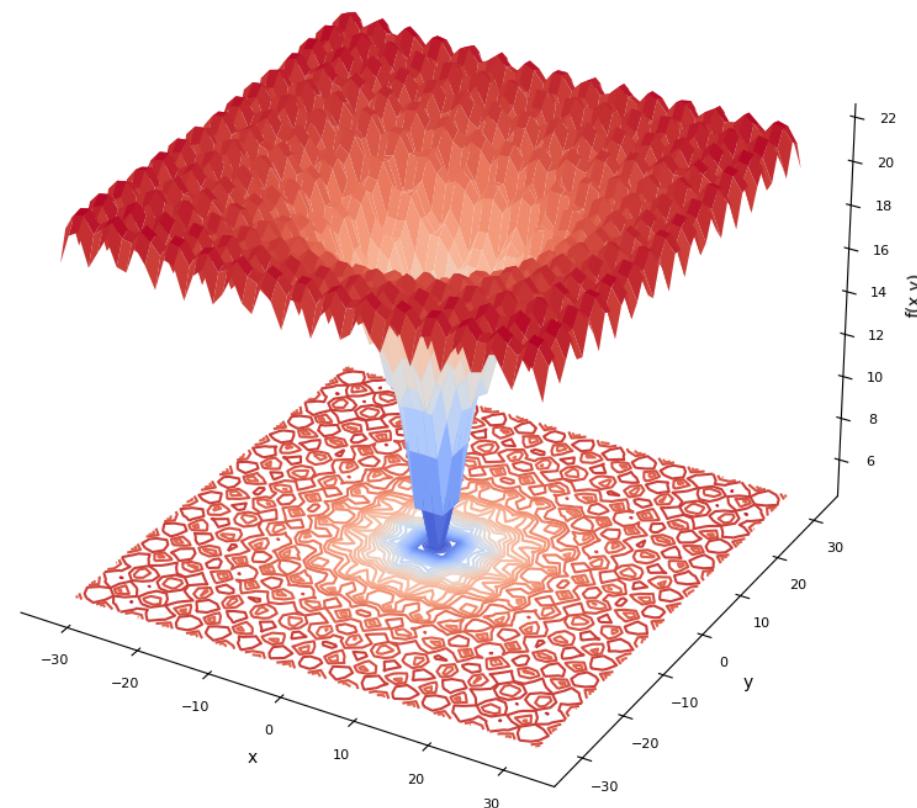
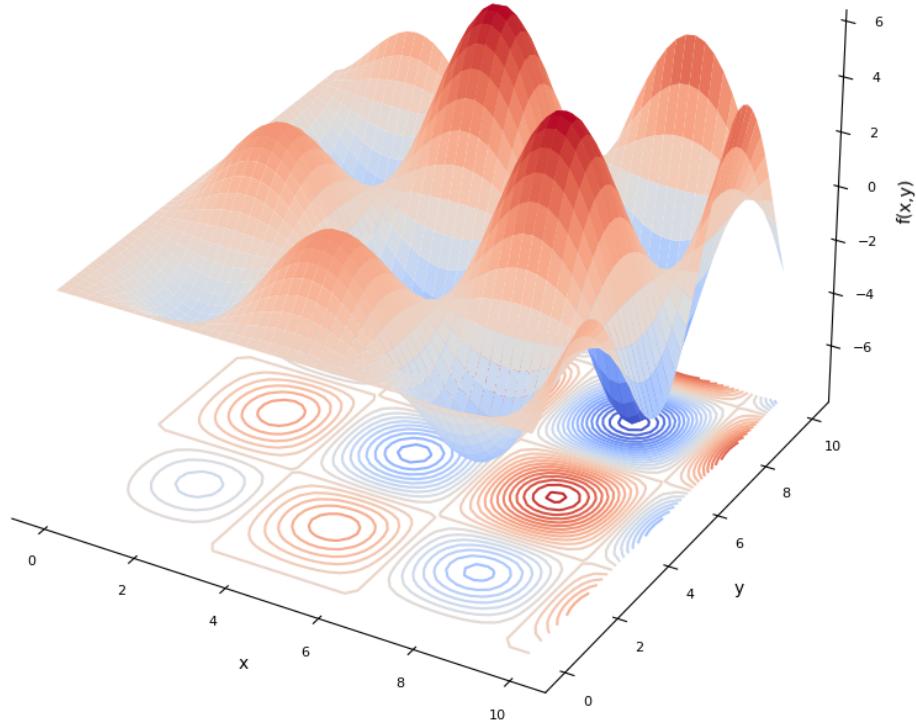
Practice Session 5: Constrained optimization (10 min)

Global optimization

Global vs local optimization

- Local: Find any local optimum
 - All we have done so far
- Global: Find best local optimum
 - Needs bounds to be well defined
 - Extremely challenging in high dimensions
- Same for convex problems

Examples



Genetic algorithms

- Heuristic inspired by natural selection
- Random initial population of parameters
- In each evolution step:
 - Evaluate "fitness" of all population members
 - Replace worst by combinations of good ones
- Converge when max iterations are reached
- Examples: "pygmo_gaco" , "pygmo_bee_colony" , "nlopt_crs2_lm" , ...

Bayesian optimization

- Evaluate criterion on grid or sample of parameters
- Build surrogate model of criterion
- In each iteration
 - Do new criterion evaluations at promising points
 - Improve surrogate model
- Converge when max iterations is reached

Multistart optimization (in estimagic)

- Inspired by [tiktak algorithm](#) by Fatih Guvenen and Serdar Ozkan
- Evaluate criterion on random exploration sample
- Run local optimization from best point
- In each iteration:
 - Combine best parameter and next best exploration point
 - Run local optimization from there
- Converge if current best optimum was rediscovered several times
- Use any estimagic algorithm for local optimization
- Can distinguish soft and hard bounds

How to choose

- Extremely expensive criterion (i.e. can only do a few evaluations):
 - Bayesian optimization
- Differentiable function or least-squares structure and not too many local optima:
 - Multistart with a local optimizer tailored to function properties
- Rugged function with extremely many local optima
 - Genetic optimizer
 - Consider refining the result with local optimizer
- All are equally parallelizable

Advanced multistart

- Many local optima:
 - Large "n_sample"
 - Low "share_optimizations"
 - Weak convergence criteria
 - Refine result with stricter convergence criteria
- Few local optima:
 - Stick with defaults

Scaling

What is scaling

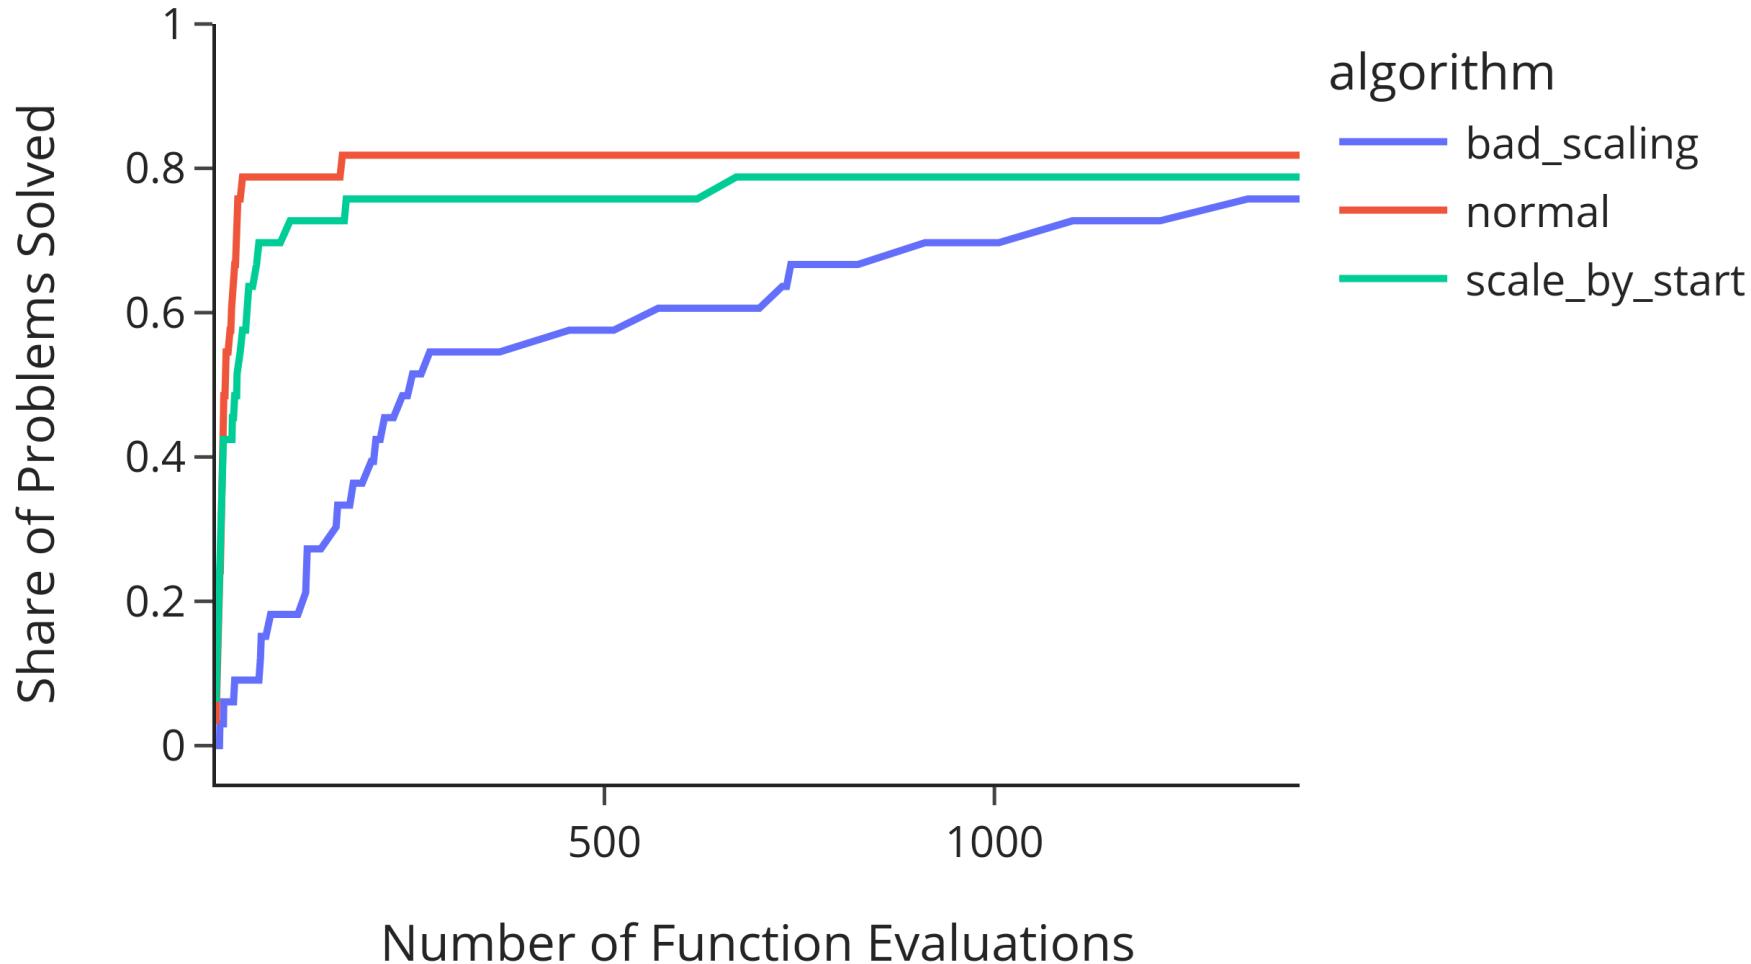
- Single most underrated topic among economists!
- Well scaled: A fixed step in any parameter dimension yields roughly comparable changes in function value
 - $f(a, b) = 0.5a^2 + 0.8b^2$
- Badly scaled: Some parameters are much more influential
 - $f(a, b) = 1000a^2 + 0.2b^2$
- Often arises when parameters have very different units

Effects of bad scaling

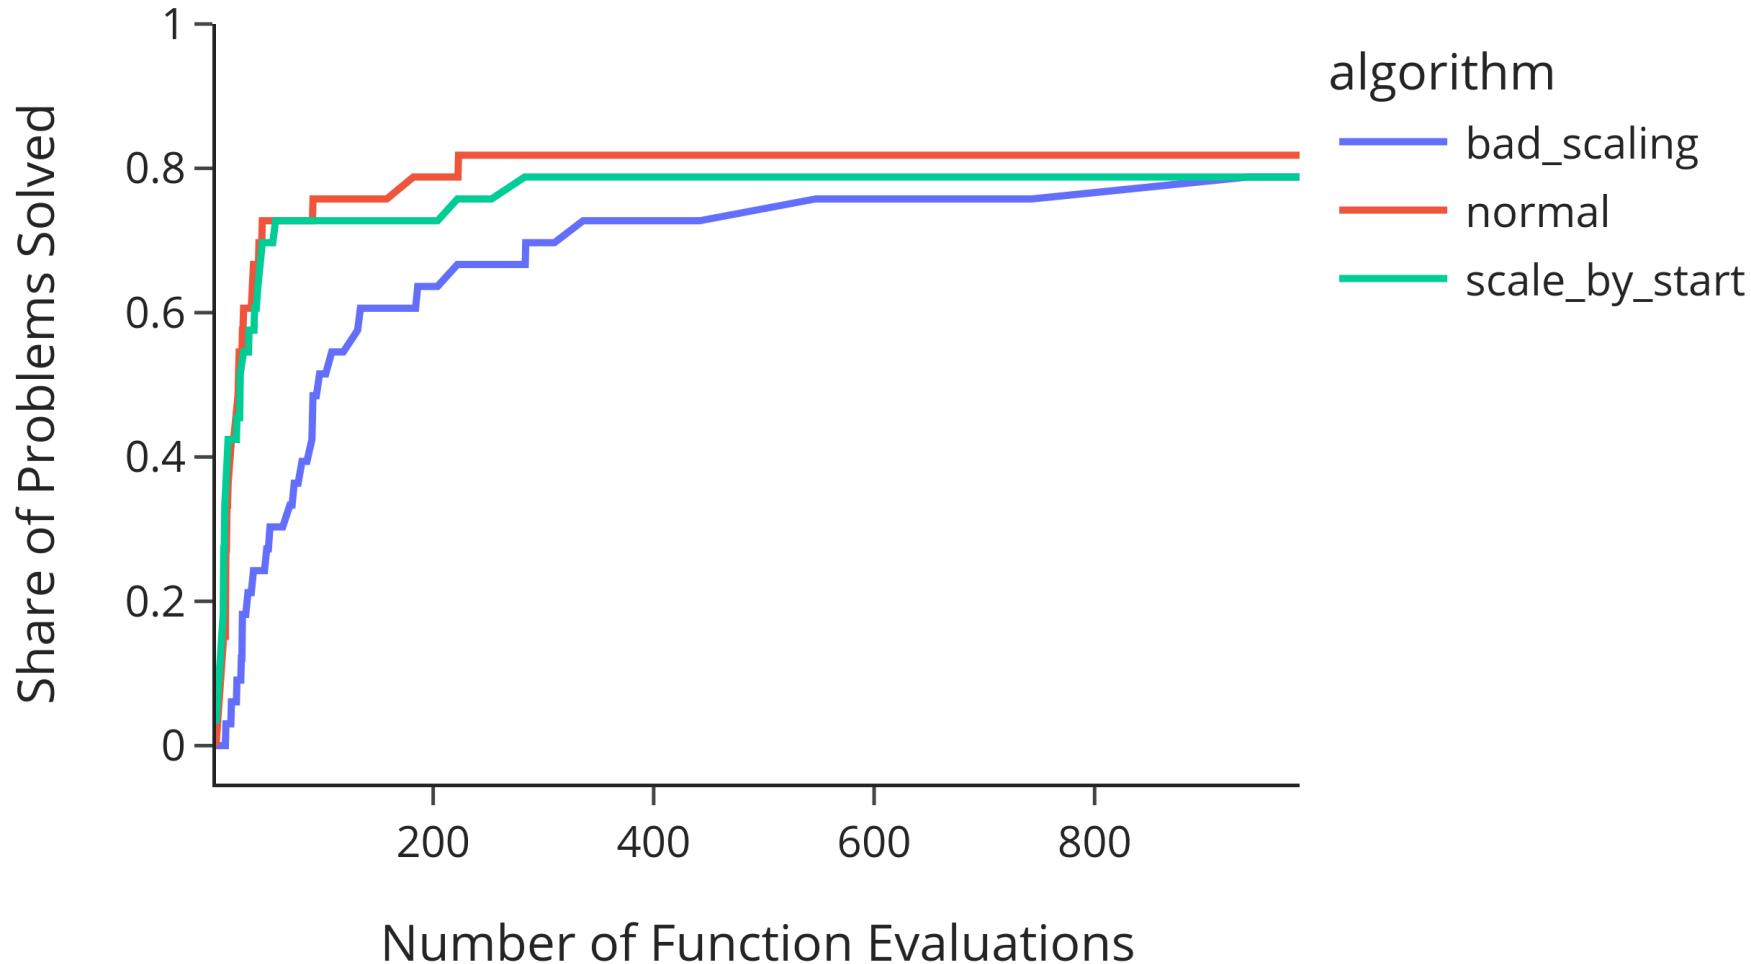
- Will try performance of optimizers on
 - Standard "estimagic" set of problems
 - badly scaled version
 - badly scaled version when `scaling=True` in `minimize`

```
problems_bad_scaling = get_benchmark_problems(  
    name="estimagic",  
    scaling=True,  
    scaling_options={"min_scale": 1, "max_scale": 1_000},  
)
```

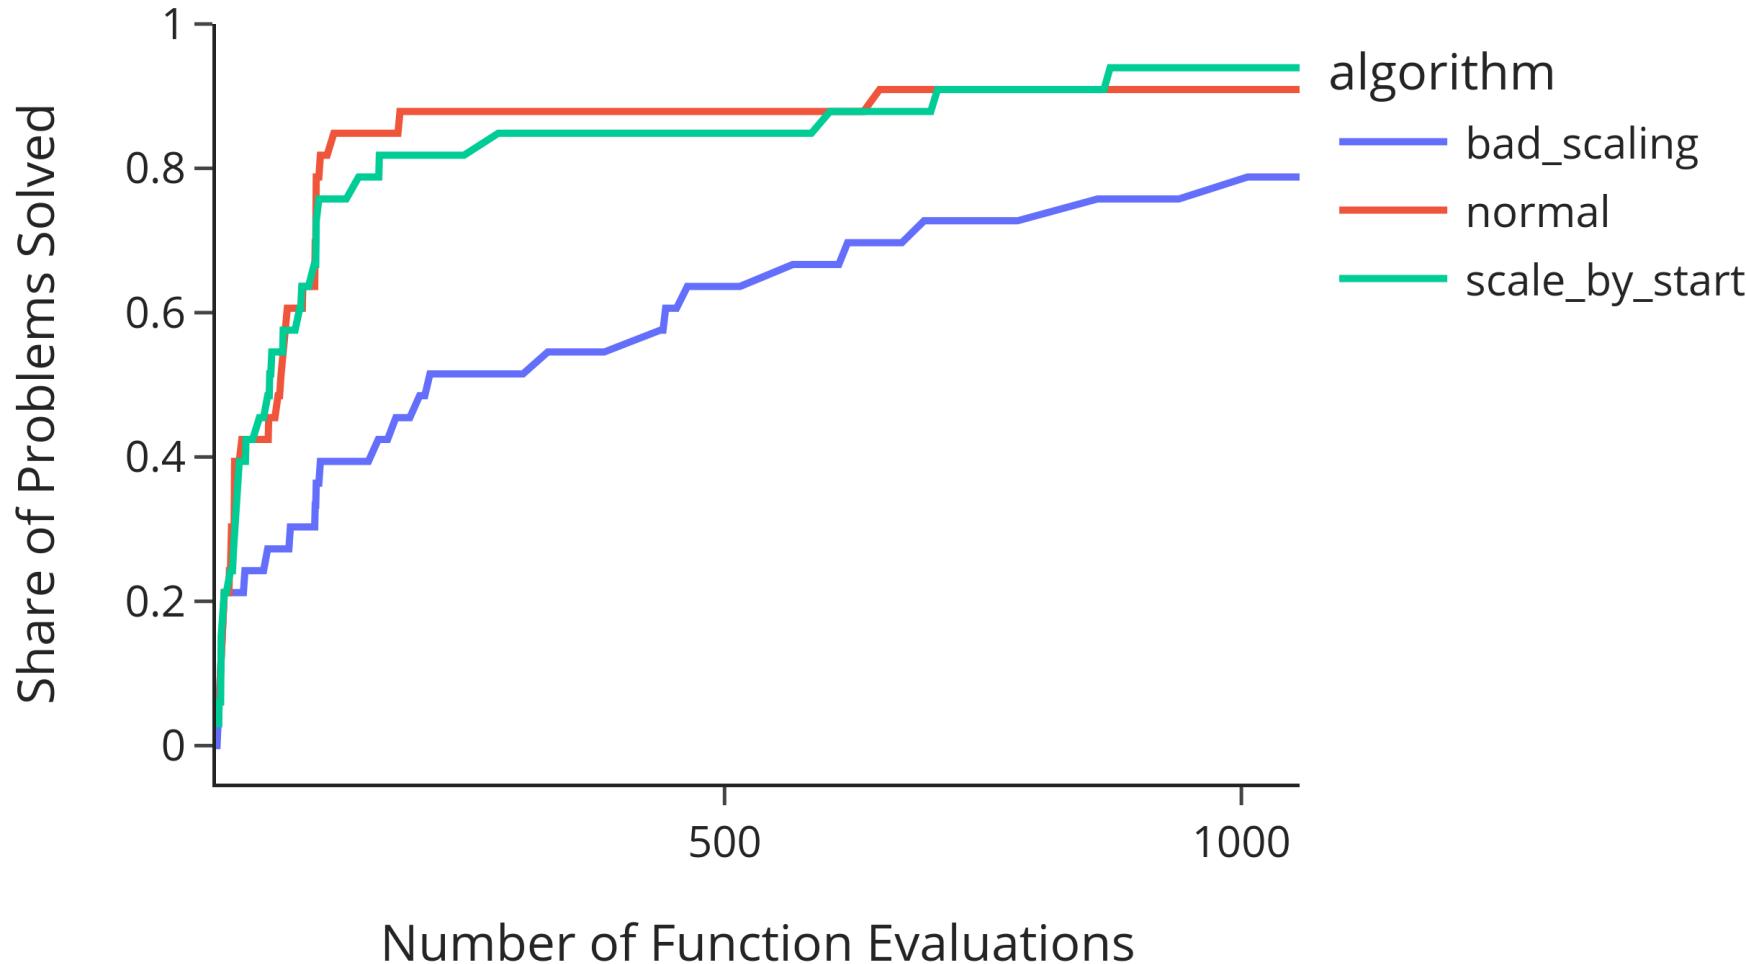
Effect of bad scaling: `scipy_lbfgsb`



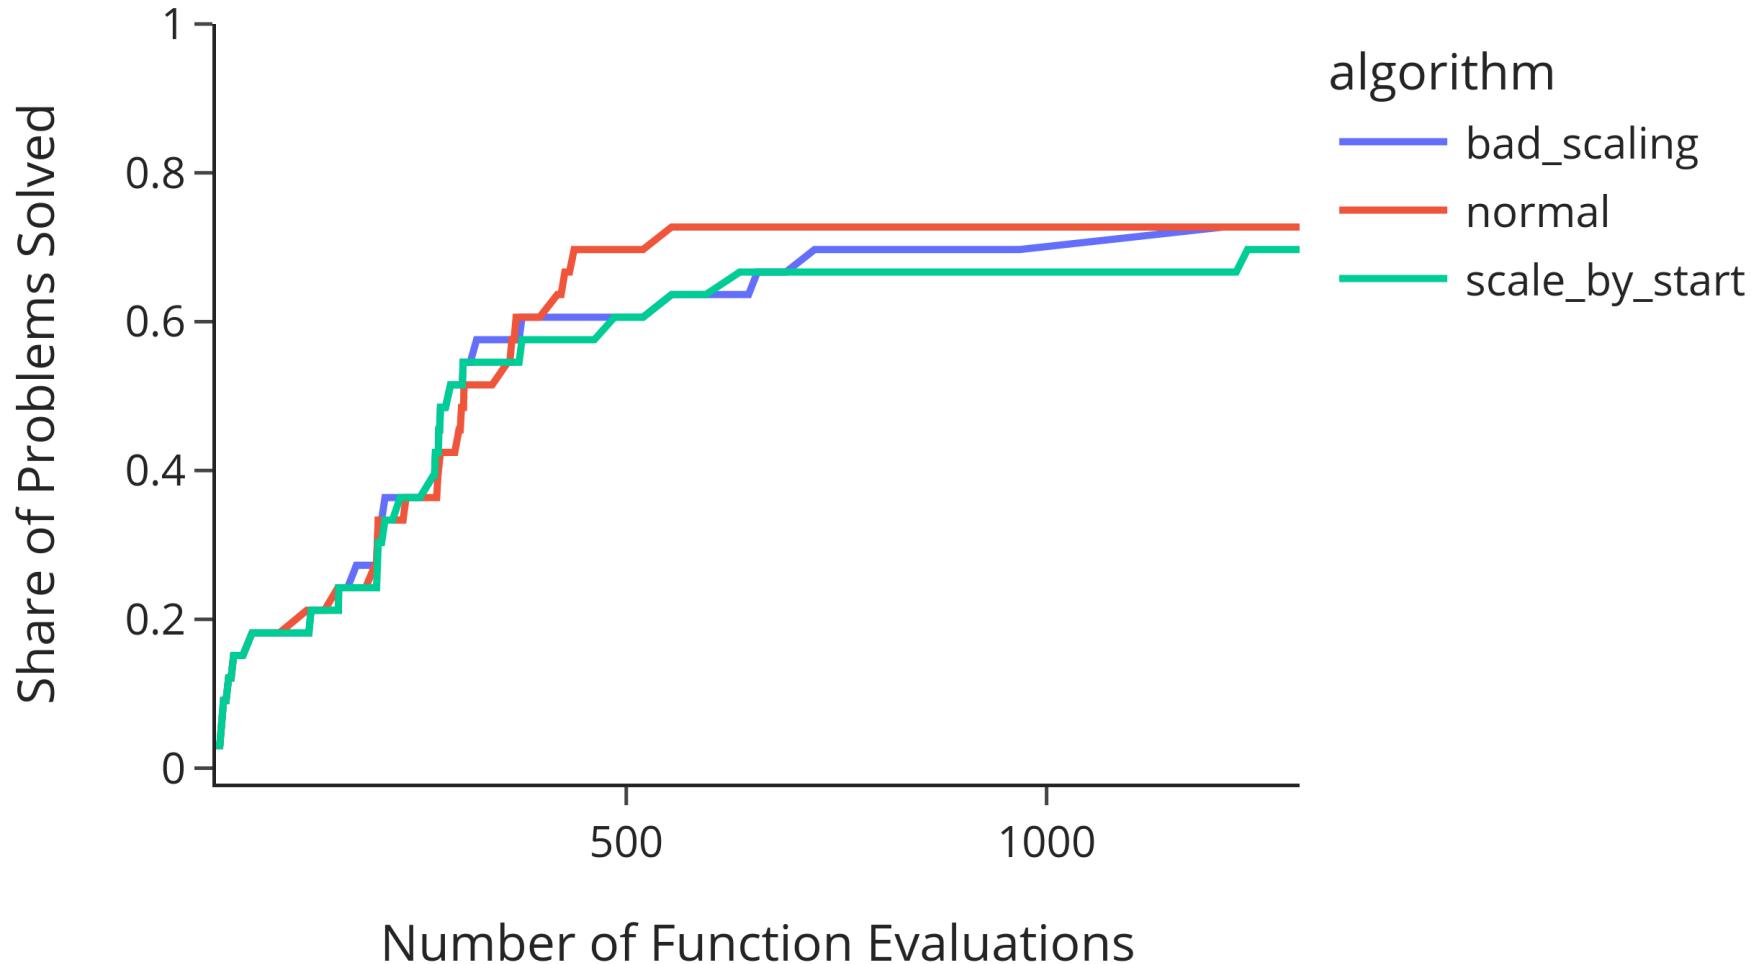
Effect of bad scaling: fides



Effect of bad scaling: nag_dfols



Effect of bad scaling: nlopt_bobyqa



Scaling in estimagic: By start params

```
def badly_scaled(x):
    out = (
        0.01 * np.abs(x[0])
        + np.abs(x[1])
        + (x[2] - 0.9) ** 6
    )
    return out

em.minimize(
    criterion=badly_scaled,
    params=np.array([200, 1, 1]),
    scaling=True,
)
```

- Optimizer sees x / x_{start}
- Works without bounds
- Will recover that $x[0]$ needs large steps
- Won't recover that $x[2]$ needs tiny steps

Scaling in estimagic: By bounds

```
em.minimize(  
    criterion=badly_scaled,  
    params=np.array([200, 1, 1]),  
    scaling=True,  
    lower_bounds=np.array([-200, -2, 0.8]),  
    upper_bounds=np.array([600, 2, 1]),  
    scaling=True,  
    scaling_options={"method": "bounds"},  
)
```

- Internal parameter space mapped to $[0, 1]^n$
- Will work great in this case
- Requires careful specification of bounds

Very scale sensitive

- nag_pybobyqa
- tao_pounders
- pounders
- nag_dfols
- scipy_cobyla

Somewhat scale sensitive

- scipy_lbfgsb
- fides

Not scale sensitive

- scipy_neldermead
- nlopt_neldermead
- nlopt_bobyqa
- scipy_powell
- scipy_ls_lm
- scipy_ls_trf

Practice Session 6: Scaling of optimization problems (10 min)

How to contribute

- Make issues or provide feedback
- Improve or extend the documentation
- Suggest, wrap or implement new optimizers
- Teach estimagic to colleagues, students and friends
- Make us happy by giving us a  on
github.com/OpenSourceEconomics/estimagic

Numerical vs automatic differentiation (oversimplified)

- **Automatic differentiation**

- Magic way to calculate precise derivatives of Python functions
- Gradient calculation takes 3 to 4 times longer than function
- Runtime is independent of number of parameters
- Code must be written in a certain way

- **Numerical differentiation**

- Finite step approximation to derivatives
- Less precise than automatic differentiation
- Runtime increases linearly with number of parameters

What is JAX

- GPU accelerated replacement for Numpy
- State of the art automatic differentiation
 - Gradients
 - Jacobians
 - Hessians
- Just in time compiler for python code
- Composable function transformations such as `vmap`
- The 2nd best Python package (after estimagic)

Calculating derivatives with JAX

```
>>> import jax.numpy as jnp
>>> import jax

>>> def sphere(x):
...     return jnp.sum(x ** 2)

>>> gradient = jax.grad(sphere)

>>> gradient(jnp.array([1., 1.5, 2.]))
DeviceArray([2., 3., 4.], dtype=float64)
```

Practice Session 7: Using JAX derivatives in estimagic (10 min)

What is JAXopt

- Library of optimizers written in JAX
- Hardware accelerated
- Batchable
- Differentiable

When to use it

- Solve many instances of same optimization problem
- Differentiate optimization w.r.t hyperparameters

Simple optimization in JAXopt

```
>>> import jax
>>> import jax.numpy as jnp
>>> from jaxopt import LBFGS

>>> x0 = jnp.array([1.0, 2, 3])
>>> shift = x0.copy()

>>> def criterion(x, shift):
...     return jnp.vdot(x, x + shift)

>>> solver = LBFGS(fun=criterion)

>>> result = solver.run(init_params=x0, shift=shift)
>>> result.params
DeviceArray([ 0. , -0.5, -1. ], dtype=float64)
```

- import solver
- initialize solver with criterion
- run solver with starting parameters
- pass additional arguments of criterion to run method

Vmap

```
>>> import numpy as np
>>> import scipy as sp

>>> a = np.stack([np.diag([1, 2]), np.diag([3, 4])])
>>> a[0]
array([[ 1.,  0.],
       [ 0.,  2.]])  
  
>>> sp.linalg.inv(a[0])
array([[ 1., -0.],
       [ 0.,  0.5]])  
  
>>> sp.linalg.inv(a)
... ValueError: expected square matrix  
  
>>> res = [sp.linalg.inv(a[i]) for i in [0, 1]]
>>> np.stack(res)
array([[[ 1.        , -0.        ],
         [ 0.        ,  0.5       ]],
        [[ 0.33333333, -0.        ],
         [ 0.        ,  0.25      ]]])
```

- consider matrix inversion
- not defined for arrays with dimension > 2 (in `scipy`)
- loop over 2d matrices
- syntactic sugar: `np.vectorize` and the like
 - does not increase speed

Vmap in JAX

```
>>> import jax.numpy as jnp
>>> import jax.scipy as jsp
>>> from jax import vmap, jit

>>> a = jnp.array(a)

>>> jax_inv = jit(vmap(jsp.linalg.inv))

>>> jax_inv(a)
DeviceArray([[ [1.        ,  0.        ],
              [0.        ,  0.5       ],
              [[0.33333333,  0.        ],
               [0.        ,  0.25      ]]], dtype=float64)
```

- use `jax.numpy` and `jax.scipy`
- define vectorized map using `jax.vmap`
- need `jit` on the outside to compile the new function

Vectorize an optimization in JAXopt

```
>>> from jax import jit, vmap

>>> def solve(x, shift):
...     return solver.run(init_params=x, shift=shift).params

>>> batch_solve = jit(vmap(solve, in_axes=(None, 0)))
>>> weights = jnp.array([
    [0.0, 1.0, 2.0],
    [3.0, 4.0, 5.0]
])
>>> batch_solve(x0, weights)
DeviceArray([[ 0. , -0.5, -1. ],
            [-1.5, -2. , -2.5]], dtype=float64)
```

- import jit and vmap
- define wrapper around solve
- call vmap on wrapper
 - in_axes=(None, 0) means that we map over the 0-axis of the second argument
 - call jit at the end

Differentiate a solution in JAXopt

```
>>> from jax import jacobian

>>> jacobian(solve, argnums=1)(x0, weight)
DeviceArray([[-0.5,  0.,  0.],
            [ 0., -0.5,  0.],
            [ 0.,  0., -0.5]], dtype=float64)

>>> solve(x0, weight)
DeviceArray([ 0., -0.5, -1.], dtype=float64)

>>> solve(x0, weight + 1)
DeviceArray([-0.5, -1., -1.5], dtype=float64)
```

- import jacobian or grad
- use argnums to specify by which argument we differentiate

Practice Session 8: Vectorized optimization in JAXopt (15 min)

Summary

Summary: Libraries

- Use scipy if you want as few dependencies as possible
- Use jaxopt for:
 - batch optimization of many small problems
 - differentiating optimizers w.r.t tuning parameters
- Use estimagic if you need:
 - Access to a huge number of optimizers
 - Advanced diagnostics
 - Advanced options such as scaling

Summary: Tips

- Exploit least squares structure if you can
- Look at criterion and params plots
- Try multiple algorithms (based on theory) and take the best (based on experiments)
- Try to improve the scaling of your problem
- Use multistart over genetic algorithms if finding local optima is easy
- Make your problem jax compatible and use automatic differentiation